

Advanced Statistics : Real estate - Boston area

[Code ▼](#)

Introduction

- L'immobilier aux US est plutôt une thématique fameuse dans son genre. Notamment au centre de la crise des subprimes, il peut être pertinent de suivre l'évolution du prix du logement aux US (bdd date des années 70)
- Problématique : Quelles variables peuvent expliquer le prix médian d'un bien immobilier dans la banlieue de Boston ?
- Pour ce faire, nous avons grâce au cours de Big Dad et au cours de R développé plusieurs outils : Pour en citer les majeurs, il s'agira d'utiliser la Principal Components Analysis et Backward Stepwise Regression mais aussi Subset Sélection, Cross Validation ou la ridge selection.
- Choix de la variable dépendante : MEDV. Nous avons fait des essais avec les différentes variables, c'est en prenant MEDV comme variables dépendante que le r squared était le plus élevé (74%). De plus, une analyse qualitative des différentes variables nous a permis de soutenir ce choix, l'impact que pouvaient avoir de nombreuses variables sur MEDV étant évident.

Libraries Installation

[Hide](#)[Hide](#)

```
install.packages("leaps") # Best Subset Selection
library(leaps)
install.packages("glmnet") # Elastic Net
library(glmnet)
```

Subset Selection Method Code

- Les modèles qui affectent une "pénalité" à l'augmentation du nombre de variables sont plus pertinents. Les autres modèles (RSS et R squared, sans pénalité) indiquent de n'enlever aucun prédicteur.
- Les modèles du r^2 ajusté, du BIC et du C_p nous donne le même résultat :

garder le modèle avec 11 variables explicatives.

- Le “significance level” des variables restantes reste proche, sauf pour le prédicteur TAX qui devient plus significatif.

Hide

Hide

```
install.packages("leaps") # Best Subset Selection
```

```
essai de l'URL 'https://cran.rstudio.com/bin/macosx/el-capitan/contrib/3.4/leaps_3.0.tgz'
Content type 'application/x-gzip' length 69196 bytes (67 KB)
=====
downloaded 67 KB
```

```
The downloaded binary packages are in
  /var/folders/ql/qw81rhln68bcj9f1nhfrqxv40000gn/T//RtmpReI2p1/
downloaded_packages
```

Hide

Hide

```
library(leaps)
#housing<-read.csv("R/Projet R/traitement_housing.csv",sep=";")
#head(housing) ; View(housing) ; print(names(housing)) ; print(dim(housing))
sum(is.na(housing)) # 0 !
```

```
[1] 0
```

Hide

Hide

```
reglin<-lm(MEDV~.,housing)
summary(reglin) # r squared of 74% + low pval(Fstat)
```

```
Call:
lm(formula = MEDV ~ ., data = housing)

Residuals:
    Min       1Q   Median       3Q      Max
-15.595  -2.730  -0.518   1.777  26.199





Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.646e+01  5.103e+00   7.144 3.28e-12 ***
CRIM        -1.080e-01  3.286e-02  -3.287 0.001087 **
ZN          4.642e-02  1.373e-02   3.382 0.000778 ***
INDUS       2.056e-02  6.150e-02   0.334 0.738288
Chase       2.687e+00  8.616e-01   3.118 0.001925 **
NOX        -1.777e+01  3.820e+00  -4.651 4.25e-06 ***
RM          3.810e+00  4.179e-01   9.116 < 2e-16 ***
AGE         6.922e-04  1.321e-02   0.052 0.958229
DIS        -1.476e+00  1.995e-01  -7.398 6.01e-13 ***
RAD         3.060e-01  6.635e-02   4.613 5.07e-06 ***
TAX        -1.233e-02  3.760e-03  -3.280 0.001112 **
PTRATIO    -9.527e-01  1.308e-01  -7.283 1.31e-12 ***
B           9.312e-03  2.686e-03   3.467 0.000573 ***
LTSAT      -5.248e-01  5.072e-02 -10.347 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom
Multiple R-squared:  0.7406,    Adjusted R-squared:  0.7338
F-statistic: 108.1 on 13 and 492 DF,  p-value: < 2.2e-16
```

Hide

Hide

```
par(mfrow=c(2,2))
plot(reglin)
```



Hide

Hide

```
regfit.full<-regsubsets(MEDV~.,housing,nvmax=13)
regsummary=summary(regfit.full)
```

- We can notice that the QQ-plot is not so satisfying. We can therefore try a regression using the log function, we can see that regressing the log of the dependant variable on the logs of the predictors brings a better result as for the QQ-Plot.

Hide

Hide

```
predictors<-log(housing$DIS+1)+log(housing$INDUS+1)+log(housing$C
RIM+1)+log(housing$ZN+1)+log(housing$Chase+1)+log(housing$NOX+1)+
log(housing$RM+1)+log(housing$AGE+1)+log(housing$RAD+1)+log(housi
ng$TAX+1)+log(housing$PTRATIO+1)+log(housing$B+1)+log(housing$LTS
AT+1)
logreg <- lm(log(MEDV+1)~predictors,housing)
summary(logreg)
```

```
Call:
lm(formula = log(MEDV + 1) ~ predictors, data = housing)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.97098	-0.20698	-0.03523	0.19437	1.09560

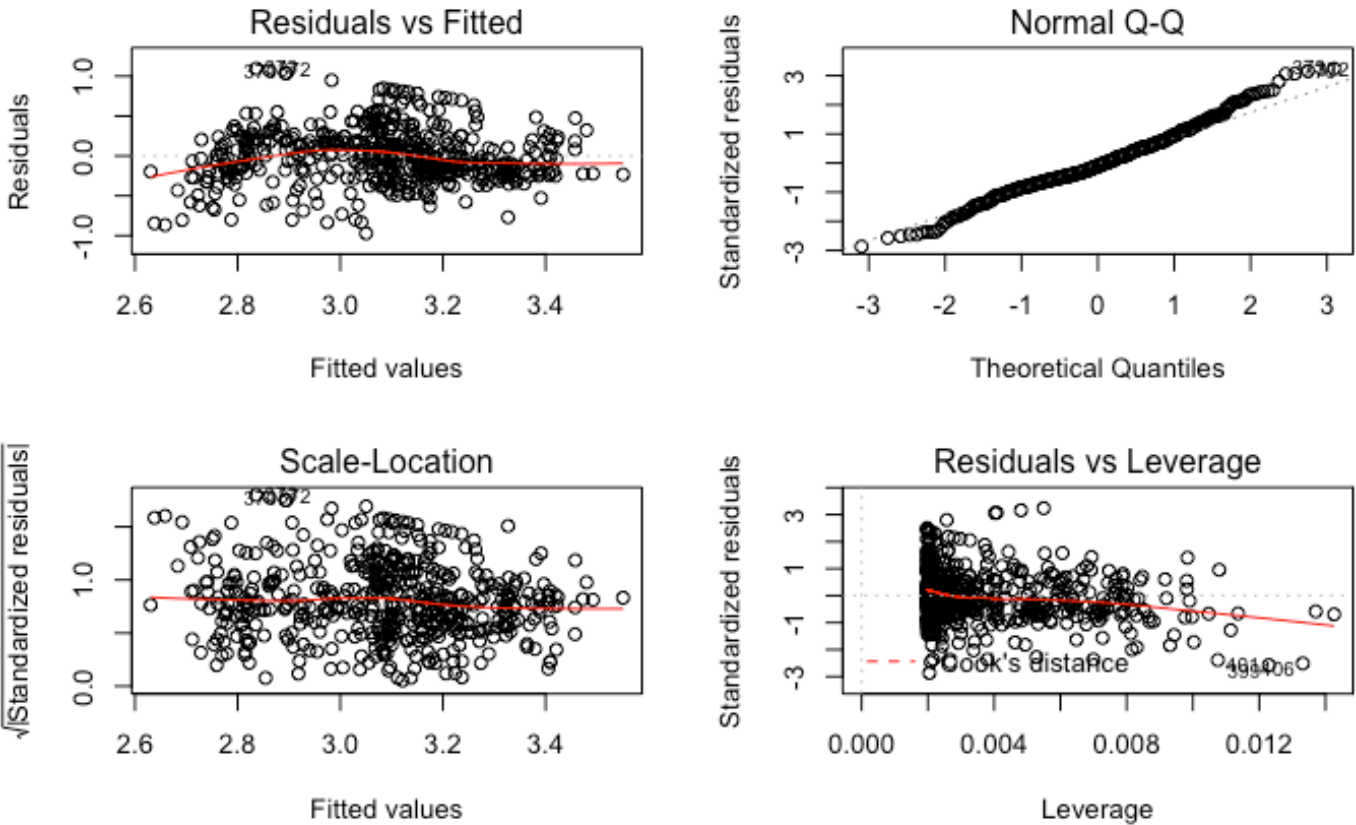
```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.541007   0.198846   27.87  <2e-16 ***
predictors  -0.078417   0.006332  -12.38  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.3392 on 504 degrees of freedom
Multiple R-squared: 0.2333, Adjusted R-squared: 0.2318
F-statistic: 153.4 on 1 and 504 DF, p-value: < 2.2e-16

Hide

Hide

```
par(mfrow=c(2,2))
plot(logreg) # better QQ-plot than with the normal regression...
```



Hide

```
# show for each subset size, the variables that should be kept to
have the best model (smallest RSS)
regsummary
```

```
Subset selection object
Call: regsubsets.formula(MEDV ~ ., housing, nvmax = 13)
13 Variables  (and intercept)
      Forced in Forced out
CRIM      FALSE      FALSE
ZN        FALSE      FALSE
INDUS     FALSE      FALSE
Chase     FALSE      FALSE
NOX       FALSE      FALSE
RM        FALSE      FALSE
AGE       FALSE      FALSE
DIS       FALSE      FALSE
RAD       FALSE      FALSE
TAX       FALSE      FALSE
PTRATIO   FALSE      FALSE
B         FALSE      FALSE
LTSAT     FALSE      FALSE
1 subsets of each size up to 13
Selection Algorithm: exhaustive
      CRIM ZN  INDUS Chase NOX RM  AGE DIS RAD TAX PTRATIO B
LTSAT
1  ( 1 )  " "  " "  " "  " "  " "  " "  " "  " "  " "  " "  " "
" "*"
2  ( 1 )  " "  " "  " "  " "  " "  "*" " "  " "  " "  " "  " "
" "*"
3  ( 1 )  " "  " "  " "  " "  " "  " "  "*" " "  " "  " "  " "*"
" "*"
4  ( 1 )  " "  " "  " "  " "  " "  " "  "*" " "  "*" " "  " "  "*"
" "*"
5  ( 1 )  " "  " "  " "  " "  " "  "*" "*" " "  "*" " "  " "  "*"
" "*"
6  ( 1 )  " "  " "  " "  "*"  "*"  "*" "*" " "  "*" " "  " "  "*"
" "*"
7  ( 1 )  " "  " "  " "  "*"  "*"  "*" " "  "*" " "  " "  "*"
" "*"
8  ( 1 )  " "  "*" " "  "*"  "*"  "*" " "  "*" " "  " "  "*"
" "*"
9  ( 1 )  "*"  " "  " "  "*"  "*"  "*" " "  "*" "*" " "  "*"
" "*

```

```
"  "*"
```

10	(1)	"*"	"*"	" "	" "	"*"	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"
		" "	"*"											
11	(1)	"*"	"*"	" "	"*"	"*"	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"
		" "	"*"											
12	(1)	"*"	"*"	"*"	"*"	"*"	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"
		" "	"*"											
13	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"
		" "	"*"											

Hide

Hide

```
names(regsummary) # different methods of selection of the best model between the different subset sizes
```

```
[1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat"
" "obj"
```

Hide

Hide

```
par(mfrow=c(3,2))
plot(regsummary$rss,xlab="nb of variables",ylab="RSS",type="l")
plot(regsummary$rsq,xlab="nb of variables",ylab="R squared",type="l")
```

Hide

Hide

```
plot(regsummary$adjr2,xlab="nb of variables",ylab="adjusted R squared",type="l")
# no big difference between r2 and adjusted r2
plot(regsummary$cp,xlab="nb of variables",ylab="Cp", type="l")
```

Hide

Hide

```
plot(regsummary$bic,xlab="number of variables",ylab="BIC",type="l")
which.min(regsummary$rss);which.max(regsummary$rsq)
```

```
[1] 13
[1] 13
```

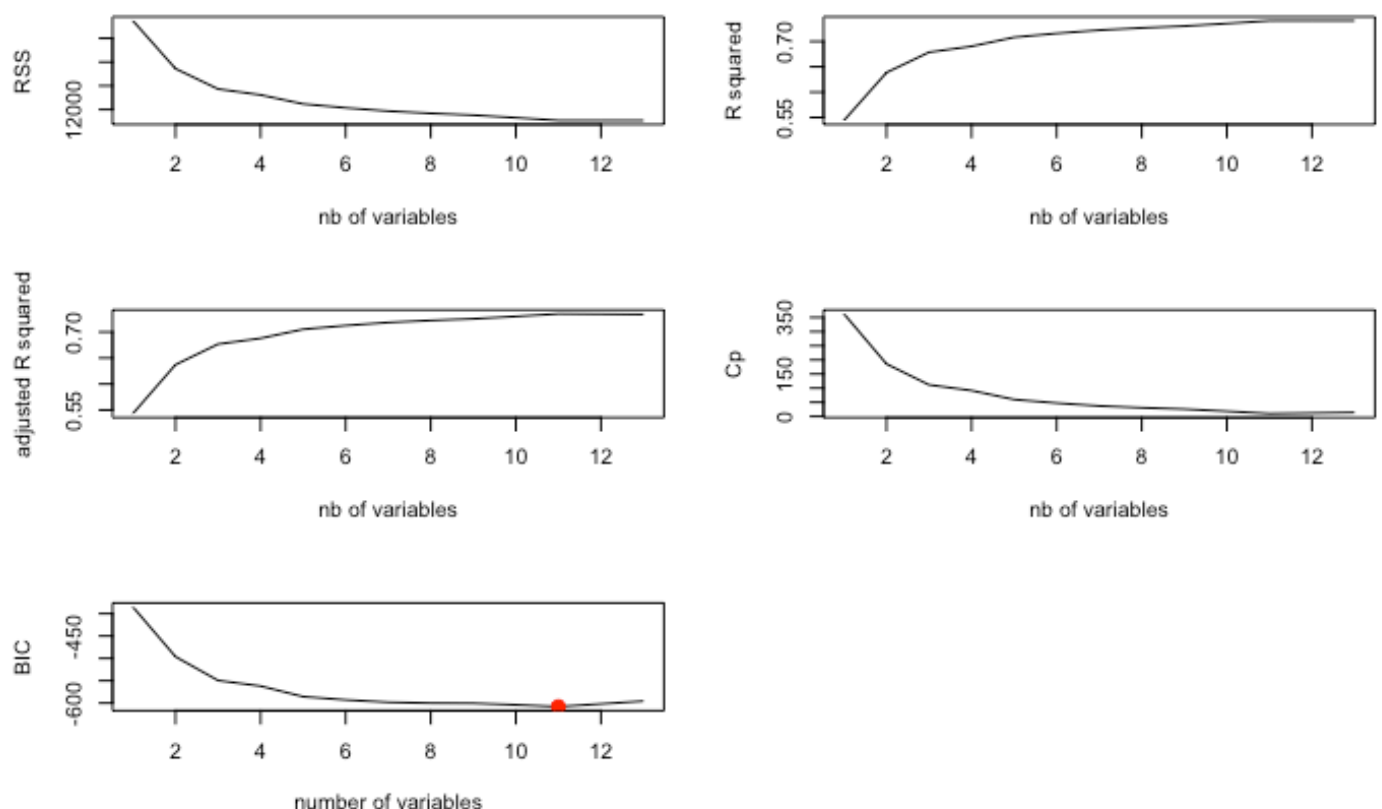
[Hide](#)[Hide](#)

```
# Same result of 13 (the number of explanatory variables) it is c
onsistent with the fact that it always increases/decreases with t
he number of variables
which.max(regsummary$adjr2);which.min(regsummary$cp);which.min(re
gsummary$bic)
```

```
[1] 11
[1] 11
[1] 11
```

[Hide](#)[Hide](#)

```
# We obtain the same result (and same as adjr2) : the model with
11 explanatory variables
points(11,regsummary$bic[11],col="red",cex=2,pch=20)
```

[Hide](#)


```
newregfit<-lm(MEDV~CRIM+ZN+Chase+NOX+RM+DIS+RAD+TAX+PTRATIO+B+LTS
AT,housing)
# new regression with the 11 remaining variables
summary(newregfit)
```

Call:

```
lm(formula = MEDV ~ CRIM + ZN + Chase + NOX + RM + DIS + RAD +
    TAX + PTRATIO + B + LTSAT, data = housing)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.5984	-2.7386	-0.5046	1.7273	26.2373

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	36.341145	5.067492	7.171	2.73e-12	***
CRIM	-0.108413	0.032779	-3.307	0.001010	**
ZN	0.045845	0.013523	3.390	0.000754	***
Chase	2.718716	0.854240	3.183	0.001551	**
NOX	-17.376023	3.535243	-4.915	1.21e-06	***
RM	3.801579	0.406316	9.356	< 2e-16	***
DIS	-1.492711	0.185731	-8.037	6.84e-15	***
RAD	0.299608	0.063402	4.726	3.00e-06	***
TAX	-0.011778	0.003372	-3.493	0.000521	***
PTRATIO	-0.946525	0.129066	-7.334	9.24e-13	***
B	0.009291	0.002674	3.475	0.000557	***
LTSAT	-0.522553	0.047424	-11.019	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.736 on 494 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7348

F-statistic: 128.2 on 11 and 494 DF, p-value: < 2.2e-16

```
summary(lm(MEDV~.,housing)) # just to compare both
```

Call:

```
lm(formula = MEDV ~ ., data = housing)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
CRIM	-1.080e-01	3.286e-02	-3.287	0.001087	**
ZN	4.642e-02	1.373e-02	3.382	0.000778	***
INDUS	2.056e-02	6.150e-02	0.334	0.738288	
Chase	2.687e+00	8.616e-01	3.118	0.001925	**
NOX	-1.777e+01	3.820e+00	-4.651	4.25e-06	***
RM	3.810e+00	4.179e-01	9.116	< 2e-16	***
AGE	6.922e-04	1.321e-02	0.052	0.958229	
DIS	-1.476e+00	1.995e-01	-7.398	6.01e-13	***
RAD	3.060e-01	6.635e-02	4.613	5.07e-06	***
TAX	-1.233e-02	3.760e-03	-3.280	0.001112	**
PTRATIO	-9.527e-01	1.308e-01	-7.283	1.31e-12	***
B	9.312e-03	2.686e-03	3.467	0.000573	***
LTSAT	-5.248e-01	5.072e-02	-10.347	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338

F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

- While reducing the number of variables, the level of significance of the remaining variables is quite the same as the model with all the predictors, except for the predictor “TAX” which has become more significant
- The multiple r^2 remains unchanged
- The adjusted r^2 slightly increases from 73,38% to 73,48% (linked to the “penalty”)

Hide

Hide

```
coef(regfit.full,11)
```

(Intercept)		CRIM	ZN	Chase	
NOX	RM	DIS			
36.341145004	-0.108413345	0.045844929	2.718716303	-17.37602	
3429	3.801578840	-1.492711460			
	RAD	TAX	PTRATIO	B	L
TSAT					
0.299608454	-0.011777973	-0.946524570	0.009290845	-0.52255	
3457					

Forward/Bacward Stepwsise Regression

- La méthode “forward” comme la méthode “backward” nous mènent au même résultat que précédemment : le modèle à 11 predictors sans “INDUS” et “AGE”.

Hide

Hide

```
install.packages("leaps")
```

```
essai de l'URL 'https://cran.rstudio.com/bin/macosx/el-capitan/co
ntrib/3.4/leaps_3.0.tgz'
Content type 'application/x-gzip' length 69196 bytes (67 KB)
=====
downloaded 67 KB
```

```
The downloaded binary packages are in
  /var/folders/ql/qw81rhln68bcj9f1nhfrqxv40000gn/T//RtmpZKJDqf/
downloaded_packages
```

Hide

Hide

```
library(leaps)
regfit.fwd<-regsubsets(MEDV~.,housing,nvmax=13,method="forward")
summary(regfit.fwd)
```

```
Subset selection object
```



```
13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
" "*" "
```

Hide

Hide

```
which.min(summary(regfit.fwd)$bic)
```

```
[1] 11
```

Hide

Hide

```
#regfit.bwd<-regsubsets(MEDV~.,housing,nvmax=13,method="backward"
)
#summary(regfit.bwd)
#which.min(summary(regfit.bwd)$bic)
# We obtain the same result as with the Subset Selection Method
```

Cross-Validation

- Pour des valeurs allant de 2 à 30 pour le nombre de séparations (nombre de “folds”), nous obtenons toujours le même résultat (modèle à 11 variables)

Hide

Hide

```
install.packages("leaps")
```

```
essai de l'URL 'https://cran.rstudio.com/bin/macosx/el-capitan/co
ntrib/3.4/leaps_3.0.tgz'
Content type 'application/x-gzip' length 69196 bytes (67 KB)
=====
downloaded 67 KB
```

```
The downloaded binary packages are in
  /var/folders/ql/qw81rhln68bcj9f1nhfrqxv40000gn/T//Rtmp8NXX3o/
downloaded_packages
```

Hide

[Hide](#)

```

library(leaps)
regfit.best=regsubsets(MEDV~.,data=housing,nvmax=13)
# ## c.v.
k=15 # ou autre ...
for(k in 1:30)
set.seed(8)
# create a vector that allocates each obs to one of the k=15 fold
s
folds=sample(1:k,nrow(housing),replace=TRUE)
cv.errors=matrix(NA,k,13, dimnames=list(NULL, paste(1:13)))
for(j in 1:k){
  best.fit=regsubsets(MEDV~.,data=housing[folds!=j,],nvmax=13) #
estimate outside the fold j
  for(i in 1:13){
    pred=predict(best.fit,housing[folds==j,],id=i)
    cv.errors[j,i]=mean((housing$MEDV[folds==j]-pred)^2)
  }
}
mean.cv.errors=apply(cv.errors,2,mean)
# mean across 13 models (k=15 folds)
mean.cv.errors

```

	1	2	3	4	5	6	7
8	9	10	11	12			
38.71550	31.80783	28.23559	28.49333	25.83182	26.78187	25.52970	25
.84712	26.42541	25.90942	23.87513	24.02937			
	13						
24.10530							

[Hide](#)[Hide](#)

```
which.min(mean.cv.errors)
```

```
11
11
```

[Hide](#)[Hide](#)


```
par(mfrow=c(1,1))
```

- Using $k=10$, the model with 11 predictors should be kept according to this method, however the different means (of the SSR values obtain in the 10 subsets) are very close, therefore we could reasonably think that using a penalty, we would obtain a lower number of predictors.

Hide

Hide

```
plot(mean.cv.errors,type='b')
```



Hide

Hide

```
reg.best=regsubsets(MEDV~.,data=housing, nvmax=13)  
coef(reg.best,11)
```

(Intercept)		CRIM		ZN		Chase	
NOX		RM		DIS			
36.341145004		-0.108413345		0.045844929		2.718716303	
3429		3.801578840		-1.492711460			
		RAD		TAX		PTRATIO	
TSAT						B	
						L	
0.299608454		-0.011777973		-0.946524570		0.009290845	
3457						-0.52255	

Ridge Regression Code

- Ridge Regression is a technique for analyzing multiple regression data that suffer from multicollinearity
- By adding a degree of bias to the regression estimates, ridge regression reduces the standard errors. It is hoped that the net effect will be to give estimates that are more reliable.
- If the number of regressors is larger than the number of data points (deg of liberty = $n-p-1 < 0$) all regresors cannot be included in the regression.

Hide

Hide

```
install.packages("glmnet")
```

```
essai de l'URL 'https://cran.rstudio.com/bin/macosx/el-capitan/co
ntrib/3.4/glmnet_2.0-16.tgz'
Content type 'application/x-gzip' length 1485951 bytes (1.4 MB)
=====
downloaded 1.4 MB
```

```
The downloaded binary packages are in
  /var/folders/ql/qw81rhln68bcj9f1nhfrqxv40000gn/T//Rtmp8NXX3o/
downloaded_packages
```

Hide

Hide

```
library(glmnet)
```



```
le package 'glmnet' a été compilé avec la version R 3.4.4Le chargement a nécessité le package : Matrix
Le chargement a nécessité le package : foreach
Loaded glmnet 2.0-16
```

Hide

Hide

```
# package for elastic net
x=model.matrix(MEDV~.,housing)[,-1]
y=housing$MEDV
grid=10^seq(10,-2,length=121)
grid
```

[1] 1.000000e+10 7.943282e+09 6.309573e+09 5.011872e+09 3.981072e+09 3.162278e+09 2.511886e+09
[8] 1.995262e+09 1.584893e+09 1.258925e+09 1.000000e+09 7.943282e+08 6.309573e+08 5.011872e+08
[15] 3.981072e+08 3.162278e+08 2.511886e+08 1.995262e+08 1.584893e+08 1.258925e+08 1.000000e+08
[22] 7.943282e+07 6.309573e+07 5.011872e+07 3.981072e+07 3.162278e+07 2.511886e+07 1.995262e+07
[29] 1.584893e+07 1.258925e+07 1.000000e+07 7.943282e+06 6.309573e+06 5.011872e+06 3.981072e+06
[36] 3.162278e+06 2.511886e+06 1.995262e+06 1.584893e+06 1.258925e+06 1.000000e+06 7.943282e+05
[43] 6.309573e+05 5.011872e+05 3.981072e+05 3.162278e+05 2.511886e+05 1.995262e+05 1.584893e+05
[50] 1.258925e+05 1.000000e+05 7.943282e+04 6.309573e+04 5.011872e+04 3.981072e+04 3.162278e+04
[57] 2.511886e+04 1.995262e+04 1.584893e+04 1.258925e+04 1.000000e+04 7.943282e+03 6.309573e+03
[64] 5.011872e+03 3.981072e+03 3.162278e+03 2.511886e+03 1.995262e+03 1.584893e+03 1.258925e+03
[71] 1.000000e+03 7.943282e+02 6.309573e+02 5.011872e+02 3.981072e+02 3.162278e+02 2.511886e+02
[78] 1.995262e+02 1.584893e+02 1.258925e+02 1.000000e+02 7.943282e+01 6.309573e+01 5.011872e+01
[85] 3.981072e+01 3.162278e+01 2.511886e+01 1.995262e+01 1.584893e+01 1.258925e+01 1.000000e+01
[92] 7.943282e+00 6.309573e+00 5.011872e+00 3.981072e+00 3.162278e+00 2.511886e+00 1.995262e+00
[99] 1.584893e+00 1.258925e+00 1.000000e+00 7.943282e-01 6.309573e-01 5.011872e-01 3.981072e-01
[106] 3.162278e-01 2.511886e-01 1.995262e-01 1.584893e-01 1.258925e-01 1.000000e-01 7.943282e-02
[113] 6.309573e-02 5.011872e-02 3.981072e-02 3.162278e-02 2.511886e-02 1.995262e-02 1.584893e-02
[120] 1.258925e-02 1.000000e-02

Hide

Hide

```
ridge.mod=glmnet(x,y,alpha=0,lambda=grid)
dim(coef(ridge.mod))
```

[1] 14 121

[Hide](#)[Hide](#)

```
ridge.mod$lambda[50]
```

```
[1] 125892.5
```

[Hide](#)[Hide](#)

```
coef(ridge.mod)[,50]
```

(Intercept)	CRIM	ZN	INDUS	C
hase	NOX	RM		
2.253476e+01	-3.029073e-05	1.036971e-05	-4.730956e-05	4.631163
e-04	-2.474185e-03	6.641561e-04		
AGE	DIS	RAD	TAX	PTR
ATIO	B	LTSAT		
-8.984429e-06	7.961041e-05	-2.940414e-05	-1.865180e-06	-1.573893
e-04	2.450736e-06	-6.931583e-05		

[Hide](#)[Hide](#)

```
sqrt(sum(coef(ridge.mod)[-1,50]^2))
```

```
[1] 0.002610996
```

[Hide](#)[Hide](#)

```
predict(ridge.mod,s=50,type="coefficients")[1:14,]
```

(Intercept)	CRIM	ZN	INDUS	Chase
NOX	RM	AGE		
23.599833103	-0.036202778	0.011669706	-0.052515834	0.904518380
-2.578718769	1.124746737	-0.008817626		
DIS	RAD	TAX	PTRATIO	B
LTSAT				
0.023153757	-0.026417452	-0.002035220	-0.236342882	0.003131382
-0.105893106				

[Hide](#)
[Hide](#)

```
# Train and validate
set.seed(121)
train=sample(1:nrow(x),nrow(x)/2) # 2 datasets of same size
test=(-train)
y.test=y[test]
# first try with lambda=4
ridge.mod=glmnet(x[train,],y[train],alpha=0,lambda=grid,thresh=1e-12) # we calibrate on the "train" set of data
ridge.pred=predict(ridge.mod,s=4,newx=x[test,]) # we predict on the "test" set using the "train" calibration to predict
mean((ridge.pred-y.test)^2) # out of sample error
```

```
[1] 31.03405
```

[Hide](#)
[Hide](#)

```
mean((mean(y[train])-y.test)^2) # error relative to mean
```

```
[1] 79.80175
```

[Hide](#)
[Hide](#)

```
# then we can try with very high lambda, as lambda increases, the
coefficient are driven toward zero, therefore the dependant variable
should be close to the mean
ridge.pred=predict(ridge.mod,s=1e10,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
[1] 79.80175
```

- As forecasted, it is equal to the precedent result. Indeed, when lambda is very high (~infinite penalty), the betas of the regression calibrated on “train” are close to 0, so the predicted variables are merely equals to the intercept, meaning the mean of the “train” dataset. In other words :
`ridge.pred=mean(y[train])`

[Hide](#)[Hide](#)

```
ridge.pred=predict(ridge.mod,s=0,newx=x,x=x,y=y,exact=T)
mean((ridge.pred-y.test)^2)
```

```
[1] 161.321
```

[Hide](#)[Hide](#)

```
lm(y~x,subset=train)
```

Call:

```
lm(formula = y ~ x, subset = train)
```

Coefficients:

(Intercept)	xCRIM	xZN	xINDUS	xChase
xNOX	xRM	xAGE		
32.135256	-0.129002	0.057459	-0.036471	3.536340
-15.429422	4.159140	-0.015128		
xDIS	xRAD	xTAX	xPTRATIO	xB
xLTSAT				
-1.570639	0.213266	-0.008480	-0.870637	0.008291
-0.467188				

[Hide](#)[Hide](#)


```
predict(ridge.mod,s=0,exact=T,x=x,y=y,type="coefficients")[1:14,]
```

(Intercept)		CRIM		ZN		INDUS		C	
hase		NOX		RM					
36.459353996		-0.108010484		0.046420321		0.020556585		2.68674	
0920		-17.766522516		3.809870036					
		AGE		DIS		RAD		TAX	
ATIO		B		LTSAT				PTR	
0.000692156		-1.475566864		0.306044516		-0.012334382		-0.95274	
4778		0.009311670		-0.524758345					

Hide

Hide

```
# With the Cross Validation Method
set.seed(121)
cv.out=cv.glmnet(x[train,],y[train],alpha=0) # 10-fold CV by default
plot(cv.out) #log !
```



Hide

Hide

```
bestlambda=cv.out$lambda.min
bestlambda # 0.79
```

```
[1] 0.7920875
```

Hide

Hide

```
ridge.pred=predict(ridge.mod,s=bestlambda,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
[1] 28.27471
```

Hide

Hide

```
out=glmnet(x,y,alpha=0)
predict(out,type="coefficients",s=bestlambda)[1:14,]
```

(Intercept)		CRIM		ZN		INDUS		C	
hase		NOX		RM					
27.212796997		-0.085780668		0.031448692		-0.041809644		2.90994	
6108	-11.319916884	4.017693673							
	AGE		DIS		RAD		TAX		PTR
ATIO		B		LTSAT					
-0.004104177		-1.076112442		0.141953232		-0.005342391		-0.84374	
8157	0.009023398	-0.465463708							

- We can notice that most betas decrease in absolute value except INDUS, CHASE, RM and AGE. However, INDUS and AGE were already very low with the OLS regression, so we could suggest that CHASE and LM have a really significant impact on MEDV. (they withstand the ridge effect ...)

PC ANALYSIS

Loading de la base de données et traitement de base

Hide

Hide

```
#housingdataset=read.csv("C:/Users/alexg/Desktop/Informatique/R/p
rojet/traitemnt_housing.csv")
regression_initiale=lm(housing$MEDV~.,housing)
summary(regression_initiale)
```

Call:

```
lm(formula = housing$MEDV ~ ., data = housing)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
CRIM	-1.080e-01	3.286e-02	-3.287	0.001087	**
ZN	4.642e-02	1.373e-02	3.382	0.000778	***
INDUS	2.056e-02	6.150e-02	0.334	0.738288	
Chase	2.687e+00	8.616e-01	3.118	0.001925	**
NOX	-1.777e+01	3.820e+00	-4.651	4.25e-06	***
RM	3.810e+00	4.179e-01	9.116	< 2e-16	***
AGE	6.922e-04	1.321e-02	0.052	0.958229	
DIS	-1.476e+00	1.995e-01	-7.398	6.01e-13	***
RAD	3.060e-01	6.635e-02	4.613	5.07e-06	***
TAX	-1.233e-02	3.760e-03	-3.280	0.001112	**
PTRATIO	-9.527e-01	1.308e-01	-7.283	1.31e-12	***
B	9.312e-03	2.686e-03	3.467	0.000573	***
LTSAT	-5.248e-01	5.072e-02	-10.347	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom





Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338

F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

Hide

Hide

```
par(mfrow=c(2,2))
plot(regression_initiale)
```

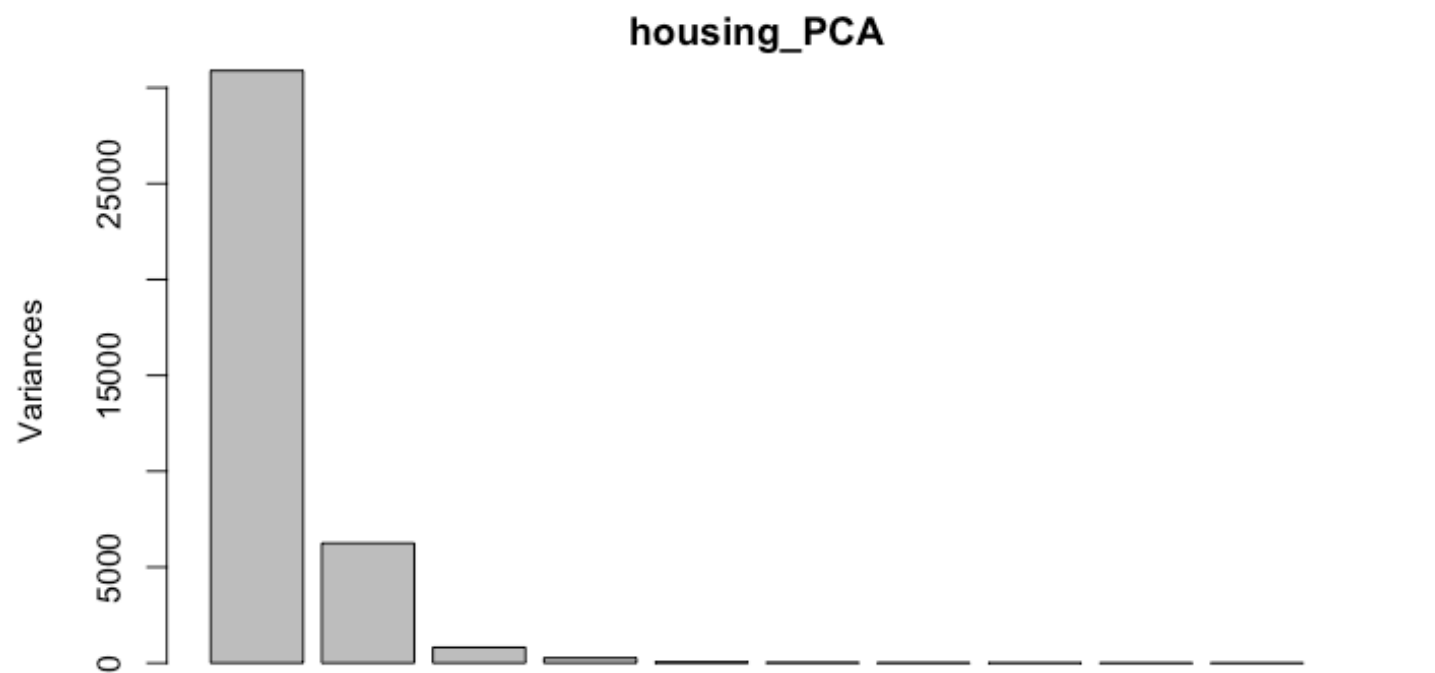



PC Creation

Hide

Hide

```
housing_PCA=prcomp(housing[, -14])
plot(housing_PCA)
```



Hide

Hide

```
summary(housing_PCA)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5
PC6	PC7	PC8	PC9	PC10	
Standard deviation		175.7553	79.0590	28.60706	16.33049
5.27985	4.00792	3.08664	1.80924	1.08671	
Proportion of Variance		0.8058	0.1631	0.02135	0.00696
0.00073	0.00042	0.00025	0.00009	0.00003	0.0013
Cumulative Proportion		0.8058	0.9689	0.99022	0.99718
0.99921	0.99963	0.99988	0.99996	0.99999	0.9985
	PC11	PC12	PC13		
Standard deviation		0.50513	0.2451	0.05527	
Proportion of Variance		0.00001	0.0000	0.00000	
Cumulative Proportion		1.00000	1.0000	1.00000	

- Par la suite,si on standardise les données afin d’éviter les problèmes de grandes différences de variances entre les variables explicatives on obtient :

Correction PC standardisées

Hide

```
housing_PCA_standadize=prcomp(housing[, -14], scale=T)
plot(housing_PCA_standadize)
```



```
summary(housing_PCA_standadize)
```

Importance of components:								
			PC1	PC2	PC3	PC4	PC5	
PC6	PC7	PC8	PC9	PC10	PC11			
Standard deviation			2.4752	1.1972	1.11473	0.92605	0.91368	0.81081
0.73168	0.62936	0.5263	0.46930	0.43129				
Proportion of Variance			0.4713	0.1103	0.09559	0.06597	0.06422	0.05057
0.04118	0.03047	0.0213	0.01694	0.01431				
Cumulative Proportion			0.4713	0.5816	0.67713	0.74310	0.80732	0.85789
0.89907	0.92954	0.9508	0.96778	0.98209				
			PC12	PC13				
Standard deviation			0.41146	0.25201				
Proportion of Variance			0.01302	0.00489				
Cumulative Proportion			0.99511	1.00000				

Corrélation entre les valeurs des PC et MEDv

Hide

Hide

```
cor(housing[,14],housing_PCA_standadize$x)
```

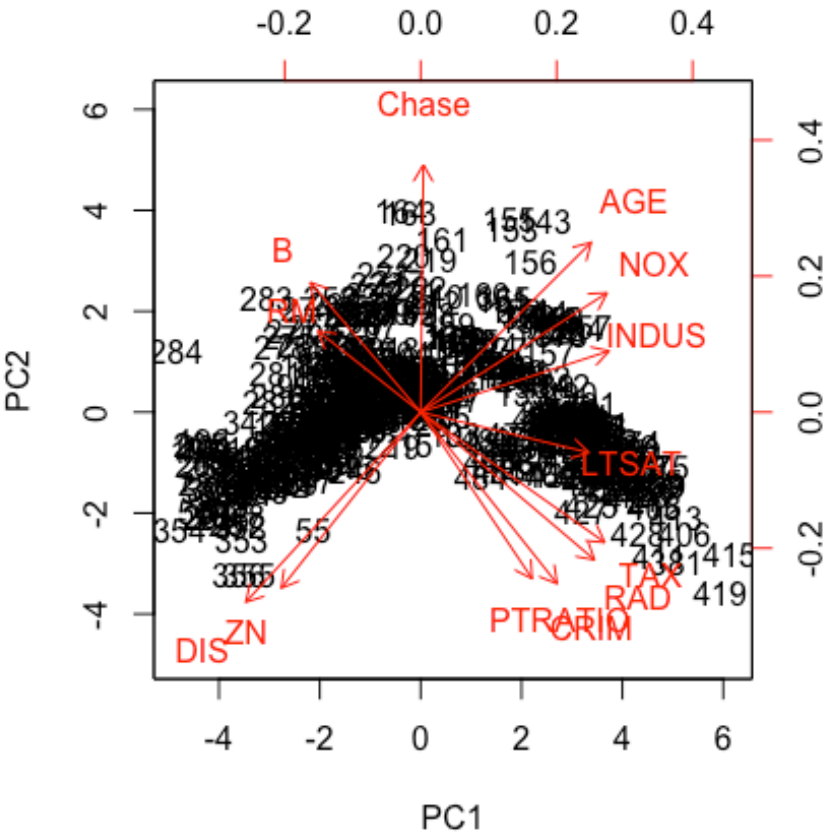
	PC1	PC2	PC3	PC4	PC5
PC6	PC7	PC8	PC9		
[1,]	-0.6117451	0.2857137	0.4243341	0.1088137	-0.2218448
219	-0.007503243	-0.07118018	0.008551394		
	PC10	PC11	PC12	PC13	
[1,]	-0.05657239	0.06441735	0.1379644	-0.09266911	

Analyse PCA et de la relation entre les variables de bases et les PC1,PC2

Hide

Hide

```
biplot(housing_PCA_standadize,scale=0)
```



- Grace au biplot, on voit comment les PC1&2 sont influencées par les variables

explicatives initiales. Notament on voit que l'age,NOX,INDUS augmentent la valeur de la PC1 de façon importante alors que DIS,ZN la diminuent et que Chase a un impact très réduit.

Construction et regression lineaire sur MEDV en fonction de toutes les PCs

Hide

Hide

```
newbase_13=cbind(housing_PCA_standadize$x,housing[ "MEDV" ] )
newbase_13
```

PC1 <dbl>	PC2 <dbl>	PC3 <dbl>	PC4 <dbl>	PC5 <dbl>
-2.0962230302	0.772348426	0.342603683	0.890892398	0.4226520
-1.4558109894	0.591399952	-0.694512011	0.486976617	-0.1956820
-2.0725465519	0.599046578	0.166956375	0.738473392	-0.9336100
-2.6089217589	-0.006863826	-0.100184990	0.343381425	-1.1038630
-2.4557547719	0.097615346	-0.075273718	0.427483833	-1.0648700
-2.2126618432	-0.009477633	-0.671716355	0.175736244	-0.6265670
-1.3575376559	0.349526292	-0.371631528	0.397740214	1.0720860
-0.8412121417	0.577228493	-0.518027787	0.537226588	1.3783240
-0.1797503956	0.342179528	-1.348304420	0.245676894	2.3479800
-1.0731221380	0.315888821	-0.557917054	0.379556648	1.4291160

1-10 of 506 rows | 1-7 of 14... Previous123456...51Next

Hide

Hide

```
regression_newbase13=lm(MEDV~.,newbase_13)
regression_newbase13
```

```
Call:
lm(formula = MEDV ~ ., data = newbase_13)
```

Coefficients:

(Intercept)	PC1	PC2	PC3	PC4
PC5	PC6	PC7		
22.53281	-2.27302	2.19491	3.50099	1.08068
-2.23308	-0.67063	-0.09431		
PC8	PC9	PC10	PC11	PC12
PC13				
-1.04018	0.14945	-1.10869	1.37366	3.08380
-3.38195				

Hide

Hide

```
summary(regression_newbase13)
```

```
Call:
lm(formula = MEDV ~ ., data = newbase_13)

Residuals:
    Min       1Q   Median       3Q      Max
-15.595  -2.730  -0.518   1.777  26.199

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  22.53281    0.21095  106.814 < 2e-16 ***
PC1          -2.27302    0.08531  -26.644 < 2e-16 ***
PC2           2.19491    0.17638   12.444 < 2e-16 ***
PC3           3.50099    0.18943   18.482 < 2e-16 ***
PC4           1.08068    0.22802    4.739 2.81e-06 ***
PC5          -2.23308    0.23111   -9.662 < 2e-16 ***
PC6          -0.67063    0.26044   -2.575  0.01031 *
PC7          -0.09431    0.28860   -0.327  0.74396
PC8          -1.04018    0.33552   -3.100  0.00204 **
PC9           0.14945    0.40126    0.372  0.70972
PC10         -1.10869    0.44996   -2.464  0.01408 *
PC11          1.37366    0.48960    2.806  0.00522 **
PC12          3.08380    0.51320    6.009 3.64e-09 ***
PC13         -3.38195    0.83791   -4.036 6.30e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom
Multiple R-squared:  0.7406,    Adjusted R-squared:  0.7338
F-statistic: 108.1 on 13 and 492 DF,  p-value: < 2.2e-16
```

- Nous venons de régresser MEDV en fonctionn de toutes les PC.
- Il s’agira en suite de déterminer quels PC garder et lesquels supprimer.

Construction et regression linéaire avec les pc de 1 à 5 (cf rule of thumb cours big dad)

- on choisit les 5 premières PC tels que la variance totale du modèle est expliquée à au moins 80%

```
newbase_5=cbind(housing_PCA_standadize$x[,1:5],housing["MEDV"])  
newbase_5
```

PC1 <dbl>	PC2 <dbl>	PC3 <dbl>	PC4 <dbl>	PC5 <dbl>
-2.0962230302	0.772348426	0.342603683	0.890892398	0.4226520
-1.4558109894	0.591399952	-0.694512011	0.486976617	-0.1956820
-2.0725465519	0.599046578	0.166956375	0.738473392	-0.9336100
-2.6089217589	-0.006863826	-0.100184990	0.343381425	-1.1038639
-2.4557547719	0.097615346	-0.075273718	0.427483833	-1.0648700
-2.2126618432	-0.009477633	-0.671716355	0.175736244	-0.6265679
-1.3575376559	0.349526292	-0.371631528	0.397740214	1.0720860
-0.8412121417	0.577228493	-0.518027787	0.537226588	1.3783240
-0.1797503956	0.342179528	-1.348304420	0.245676894	2.3479800
-1.0731221380	0.315888821	-0.557917054	0.379556648	1.4291160
1-10 of 506 rows				
Previous 1 2 3 4 5 6 ... 51 Next				

Hide

Hide

```
regression_newbase5=lm(MEDV~PC1+PC2+PC3+PC4+PC5,newbase_5)  
regression_newbase5
```

Call:
lm(formula = MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5, data = newbase_5)

Coefficients:
(Intercept) PC1 PC2 PC3 PC4
PC5
 22.533 -2.273 2.195 3.501 1.081
 -2.233

Hide


```
summary(regression_newbase5)
```

Call:

```
lm(formula = MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5, data = newbase_5
)
```

Residuals:

Min	1Q	Median	3Q	Max
-19.761	-2.893	-0.758	1.728	33.098

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	22.53281	0.22619	99.619	< 2e-16	***
PC1	-2.27302	0.09147	-24.850	< 2e-16	***
PC2	2.19491	0.18912	11.606	< 2e-16	***
PC3	3.50099	0.20311	17.237	< 2e-16	***
PC4	1.08068	0.24449	4.420	1.21e-05	***
PC5	-2.23308	0.24780	-9.012	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.088 on 500 degrees of freedom

Multiple R-squared: 0.697, Adjusted R-squared: 0.6939

F-statistic: 230 on 5 and 500 DF, p-value: < 2.2e-16

- Avec seulement les 5 premières variables on a réussi à obtenir un $R^2=0.7$ et toutes les variables sont significatives.

Corrélation entre les PCs dont la pvalue est inférieure à 1 %

```
newbase_4=cbind(housing_PCA_standadize$x[,c(1,2,3,4,5,8,11,12,13)
],housing["MEDV"])
cor(newbase_4[,-10],housing[,-14])
```

	CRIM		ZN		INDUS	Chase	NO
X	RM		AGE		DIS		
PC1	0.62116675	-0.63444186	0.858099068	0.012481272	0.8486442		
4	-0.46842215	0.776412273	-0.795900561				
PC2	-0.37741846	-0.38706849	0.134675912	0.544519026	0.2623239		
5	0.17877893	0.373498139	-0.417904755				
PC3	0.27485437	0.32980075	-0.017775345	0.323026549	0.1348419		
8	0.66210465	-0.019702590	-0.055442375				
PC4	-0.05720298	-0.11919382	-0.015877849	-0.755605366	0.1187442		
7	0.25984306	0.162250162	-0.199505129				
PC5	0.07506535	0.29294217	-0.007136952	0.079061821	0.1250407		
2	-0.38689634	0.015250137	0.090082023				
PC8	-0.09651306	0.25343194	-0.109466121	0.015521434	-0.0504248		
9	0.20564566	0.378135489	0.076663909				
PC11	-0.04728893	0.11332517	-0.130755103	0.006006486	0.0480111		
4	0.02292825	-0.198032612	-0.300047914				
PC12	-0.03569909	0.02938896	0.046577621	0.001638732	-0.3309501		
2	-0.06290175	0.087204155	-0.160858368				
PC13	0.01158046	-0.02039243	-0.063273909	0.009052647	0.0109953		
3	0.01148338	-0.009715174	-0.004611423				
	RAD		TAX		PTRATIO	B	LTS
AT							
PC1	0.79156616	0.83779482	0.507282755	-0.502407394	0.7667321		
84							
PC2	-0.32506342	-0.28667263	-0.366218209	0.285602099	-0.0889779		
36							
PC3	0.32021077	0.24606986	-0.360554349	-0.334580792	-0.2976324		
32							
PC4	-0.12256314	-0.09569383	-0.261723073	-0.156038625	-0.0642814		
59							
PC5	-0.18651151	-0.11919955	-0.533592691	-0.315775062	0.3605036		
47							
PC8	-0.05057452	-0.05209508	0.200064436	0.003098299	0.2670718		
72							
PC11	0.01576115	-0.04521501	0.075263074	0.008313145	0.1170455		
69							
PC12	0.04403735	0.08854349	-0.086242461	-0.017167595	-0.0227235		
18							
PC13	-0.15964602	0.18150635	0.005896553	-0.001124741	0.0061570		
38							

Suppression des PC7,PC6,PC10 et PC9 du premier test en fonction de leur pvalue puis regression linéaire

Hide

Hide

```
newbase_4=cbind(housing_PCA_standadize$x,housing["MEDV"] )
newbase_4
```

PC1<dbl>	PC2<dbl>	PC3<dbl>	PC4<dbl>	PC5<dbl>
-2.0962230302	0.772348426	0.342603683	0.890892398	0.4226520
-1.4558109894	0.591399952	-0.694512011	0.486976617	-0.1956820
-2.0725465519	0.599046578	0.166956375	0.738473392	-0.933610
-2.6089217589	-0.006863826	-0.100184990	0.343381425	-1.103863
-2.4557547719	0.097615346	-0.075273718	0.427483833	-1.064870
-2.2126618432	-0.009477633	-0.671716355	0.175736244	-0.626567
-1.3575376559	0.349526292	-0.371631528	0.397740214	1.072086
-0.8412121417	0.577228493	-0.518027787	0.537226588	1.378324
-0.1797503956	0.342179528	-1.348304420	0.245676894	2.347980
-1.0731221380	0.315888821	-0.557917054	0.379556648	1.429116

1-10 of 506 rows | 1-7 of 14... Previous123456...51Next

```
regression_newbase4=lm(MEDV~PC1+PC2+PC3+PC4+PC5+PC8+PC11+PC12+PC13,newbase_4)
regression_newbase4
```

Hide

Hide

```
Call:
lm(formula = MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC8 + PC11 +
    PC12 + PC13, data = newbase_4)
```

Coefficients:

(Intercept)	PC1	PC2	PC3	PC4
PC5	PC8	PC11		
22.533	-2.273	2.195	3.501	1.081
-2.233	-1.040	1.374		
PC12	PC13			
3.084	-3.382			

Hide

Hide

```
summary(regression_newbase4)
```

```
Call:
lm(formula = MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC8 + PC11 +
    PC12 + PC13, data = newbase_4)
```

```
Residuals:
      Min       1Q   Median       3Q      Max
-16.5289  -2.7838  -0.7749   1.7976  28.9197
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  22.53281    0.21285  105.863 < 2e-16 ***
PC1          -2.27302    0.08608  -26.407 < 2e-16 ***
PC2           2.19491    0.17797   12.333 < 2e-16 ***
PC3           3.50099    0.19113   18.317 < 2e-16 ***
PC4           1.08068    0.23007    4.697 3.42e-06 ***
PC5          -2.23308    0.23319   -9.576 < 2e-16 ***
PC8          -1.04018    0.33853   -3.073  0.00224 **
PC11          1.37366    0.49400    2.781  0.00563 **
PC12          3.08380    0.51781    5.955 4.91e-09 ***
PC13         -3.38195    0.84544   -4.000 7.29e-05 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.788 on 496 degrees of freedom
Multiple R-squared:  0.7338,    Adjusted R-squared:  0.729
F-statistic: 151.9 on 9 and 496 DF,  p-value: < 2.2e-16
```

- On selectionne toutes les PCs et on retire celles dont la PV value est supérieure à 1%

Comparons la significativité du test 5 et 4

Hide

Hide

```
anova(regression_newbase4, regression_newbase5)
```

Analysis of Variance Table

```
Model 1: MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5 + PC8 + PC11 + PC12 + PC13
Model 2: MEDV ~ PC1 + PC2 + PC3 + PC4 + PC5
  Res.Df  RSS Df Sum of Sq    F    Pr(>F)
1     496 11370
2     500 12944 -4   -1573.6 17.161 3.411e-13 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Il y a une différence statistique significative entre les deux modèles. On va choisir le modèle 1 qui est plus précis.

On reconstruit la predicted value en revenant sur les variables initiales(Avec le modèle sans les PC dont la pvalue > 1%)

Hide

Hide

```
beta=housing_PCA_standadize$rotation[,c(1,2,3,4,5,8,11,12,13)] %*
% regression_newbase4$coefficients[-1]
beta
```

	[,1]
CRIM	-1.0634251
ZN	0.4897006
INDUS	0.3646888
Chase	0.8106361
NOX	-2.4335788
RM	3.1938059
AGE	-0.4097755
DIS	-2.8884653
RAD	2.6018950
TAX	-2.1725453
PTRATIO	-2.0873720
B	0.4316314
LTSAT	-3.0796360

Conclusion

Le but de ce rapport était de déterminer les meilleures variables influant sur le prix du logement dans la banlieue de Boston. . Ainsi, notre étude a permis de dégager les résultats suivants :

- Le prix des logements est supérieur dans les zones avec le taux de criminalité le moins élevé .
- Un autre résultat plutôt intéressant concerne le niveau de monoxyde d'azote et la distance jusqu'au principaux centres d'emplois.
- D'une part, les gens veulent vivre proches de l'endroit où ils travaillent.
- Cependant, d'autre part, il est raisonnable de suggérer que le niveau de pollution augmente quand on se rapproche de ces grosses zones d'emplois.
- Les coefficients montrent que la distance au travail réduit plus le prix du logement que le niveau NOX. Autrement dit, quand on parle de pollution, les gens sont très sensible à la question. Cependant, ils donnent plus de valeur à un logement proche de leur zone d'emploi et donc avec un certain taux de NOX plutôt qu'un emploi plus loin mais avec un niveau NOX plus faible.
- Depuis, il n'y a aucun doute sur le fait que le niveau de pollution a augmenté et il serait intéressant d'examiner les façons dont cette dernière donnée affecte le prix du logement dans la banlieue de Boston.
- Notons tout de même que les américains sont beaucoup moins sensibles à la notion de pollution que la notre en Europe. Mais que cependant, ils deviennent de plus en plus sensibles car le sujet fait débat depuis de nombreuses années, exemple : COP21